

# 1 Part A: Writing a Cache Simulator

**Solution** The source code of data cache simulator with least recently used (LRU) cache level policy is presented on listing 1.

```
1  #define _GNU_SOURCE
2  #include <ctype.h>
3  #include <stdio.h>
4  #include <stdlib.h>
5  #include <unistd.h>
6  #include <getopt.h>
7  #include "cachelab.h"
8
9  const char *options = "h v s: E: b: t:";
10 const char *help = "Usage: ./csim-ref [-hv] -s <num> -E <num> -b <num> -t
    <file>\n
11 Options:\n
12 -h          Print this help message.\n
13 -v          Optional verbose flag.\n
14 -s <num>    Number of set index bits.\n
15 -E <num>    Number of lines per set.\n
16 -b <num>    Number of block offset bits.\n
17 -t <file>   Trace file.\n
18 \n
19 Examples:\n
20 linux> ./csim-ref -s 4 -E 1 -b 4 -t traces/yi.trace\n
21 linux> ./csim-ref -v -s 8 -E 2 -b 4 -t traces/yi.trace\n";
22
23 typedef struct cacheLineNode
24 {
25     int tag;
26     struct cacheLineNode *next;
27 } CacheLineNode;
28
29 typedef struct set
30 {
31     int currentSize;
32     int maxSize;
33     CacheLineNode *head;
34     CacheLineNode *tail;
35 } Set;
36
37 int hits;
38 int misses;
39 int evictions;
40 int blockOffsetMask;
41 int setNumberMask;
42 int dCacheSetBits;
43 int dCacheAssociativity;
44 int dCacheBlockBits;
45 Set *dCacheSets;
46 int dCacheVerboseMode;
47
48 int isFull(const Set *set)
49 {
50     return set->currentSize == set->maxSize;
51 }
52
53 void evict(Set *set)
54 {
55     CacheLineNode *evicted = set->head;
56     if (set->head->next == NULL)
57     {
58         set->head = NULL;
59         set->tail = NULL;
60     }
61     else
62     {
63         set->head = set->head->next;
64     }
65     set->currentSize--;
66
67
```

Listing 1 (Cont.): csim.c

```

68     free(evicted);
69
70     evictions++;
71     if (dCacheVerboseMode)
72         fprintf(stdout, "eviction ");
73 }
74
75 void load(Set *set, int tag)
76 {
77     if (isFull(set))
78         evict(set);
79
80     CacheLineNode *newLine = (CacheLineNode *)malloc(sizeof(CacheLineNode));
81     newLine->tag = tag;
82     newLine->next = NULL;
83
84     if (set->tail == NULL)
85     {
86         set->head = newLine;
87         set->tail = newLine;
88     }
89     else
90     {
91         set->tail->next = newLine;
92         set->tail = newLine;
93     }
94
95     set->currentSize++;
96 }
97
98 int isCached(Set *set, int tag)
99 {
100     int hit = 0;
101     CacheLineNode *prev = NULL;
102     CacheLineNode *current = set->head;
103     while (current != NULL)
104     {
105         if (current->tag == tag)
106         {
107             hit = 1;
108             if (current->next == NULL)
109             {
110                 break;
111             }
112
113             if (prev != NULL)
114             {
115                 prev->next = current->next;
116             }
117             else
118             {
119                 set->head = current->next;
120             }
121
122             set->tail->next = current;
123             set->tail = current;
124             current->next = NULL;
125             break;
126         }
127
128         prev = current;
129         current = current->next;
130     }
131
132     if (hit)
133     {
134         hits++;
135         if (dCacheVerboseMode)
136             fprintf(stdout, "hit ");
137     }
138     else
139     {
140         misses++;

```

Listing 1 (Cont.): csim.c

```

141         if (dCacheVerboseMode)
142             fprintf(stdout, "miss ");
143     }
144
145     return hit;
146 }
147
148 int createRightBitMask(int numberOfBits)
149 {
150     int mask = 0;
151     while (numberOfBits)
152     {
153         mask ^= 1 << --numberOfBits;
154     }
155
156     return mask;
157 }
158
159 void initializeDCache(int setBits, int associativity, int blockBits, int
    verboseMode)
160 {
161     hits = 0;
162     misses = 0;
163     evictions = 0;
164     dCacheSetBits = setBits;
165     dCacheAssociativity = associativity;
166     dCacheBlockBits = blockBits;
167     dCacheVerboseMode = verboseMode > 0 ? verboseMode : 0;
168     setNumberMask = createRightBitMask(dCacheSetBits);
169     blockOffsetMask = createRightBitMask(dCacheBlockBits);
170
171     int setNumber = 1 << dCacheSetBits;
172     dCacheSets = (Set *)malloc(setNumber * sizeof(Set));
173     for (int i = 0; i < setNumber; i++)
174     {
175         dCacheSets[i].currentSize = 0;
176         dCacheSets[i].maxSize = dCacheAssociativity;
177         dCacheSets[i].head = NULL;
178         dCacheSets[i].tail = NULL;
179     }
180 }
181
182 void parseAddress(unsigned int address, int *setNumber, int *tag, int
    *blockOffset)
183 {
184     *blockOffset = address & blockOffsetMask;
185     unsigned int cutOffset = address >> dCacheBlockBits;
186     *setNumber = cutOffset & setNumberMask;
187     *tag = cutOffset >> dCacheSetBits;
188 }
189
190 void dCacheOperate(Set *set, int tag)
191 {
192     if (!isCached(set, tag))
193     {
194         load(set, tag);
195     }
196 }
197
198 int dCacheSimulate(char command, int address, int bytes)
199 {
200     if (dCacheVerboseMode)
201     {
202         fprintf(stdout, "%c %x,%d ", command, address, bytes);
203     }
204
205     int setNumber;
206     int tag;
207     int offsetNumber;
208     parseAddress(address, &setNumber, &tag, &offsetNumber);
209     Set *set = &dCacheSets[setNumber];
210     switch (command)
211     {

```

Listing 1 (Cont.): csim.c

```

212     case 'S':
213     case 'L':
214     {
215         dCacheOperate(set, tag);
216         break;
217     }
218     case 'M':
219     {
220         dCacheOperate(set, tag);
221         dCacheOperate(set, tag);
222         break;
223     }
224     default:
225         return -1;
226 }
227
228 if (dCacheVerboseMode)
229     fprintf(stdout, "\n");
230
231 return 0;
232 }
233
234 int parseIntegerOption(char option, const char *optionValue, int *value)
235 {
236     if (sscanf(optionValue, "%d", value) == 0)
237     {
238         fprintf(stderr,
239             "Invalid value %s for option -%c, integer required.\n",
240             optionValue,
241             option);
242
243         return 0;
244     }
245
246     return 1;
247 }
248
249 int main(int argc, char **argv)
250 {
251     opterr = 0;
252     int optionVerboseMode = 0;
253     int optionSetBits = 0;
254     int optionAssociativity = 0;
255     int optionBlockBits = 0;
256     char *optionTraceFile = NULL;
257     int c;
258     while ((c = getopt(argc, argv, options)) != -1)
259         switch (c)
260         {
261             case 'h':
262                 fprintf(stderr, "%s", help);
263                 return 0;
264             case 'v':
265                 optionVerboseMode = 1;
266                 break;
267             case 's':
268                 if (!parseIntegerOption(c, optarg, &optionSetBits))
269                     return -1;
270                 break;
271             case 'E':
272                 if (!parseIntegerOption(c, optarg, &optionAssociativity))
273                     return -1;
274                 break;
275             case 'b':
276                 if (!parseIntegerOption(c, optarg, &optionBlockBits))
277                     return -1;
278                 break;
279             case 't':
280                 optionTraceFile = optarg;
281                 break;
282             case '?':
283                 if (optopt == 's' ||
284                     optopt == 'E' ||

```

Listing 1 (Cont.): csim.c

```

285         optopt == 'b' ||
286         optopt == 't')
287         fprintf(stderr, "Option -%c requires an argument.\n", optopt);
288     else if (isprint(optopt))
289         fprintf(stderr, "Unknown option -%c.\n", optopt);
290     else
291         fprintf(stderr, "Unknown option character \\x%x.\n", optopt);
292     return -1;
293 default:
294     abort();
295 }
296
297 if (optionSetBits == 0 ||
298     optionAssociativity == 0 ||
299     optionBlockBits == 0 ||
300     optionTraceFile == NULL)
301 {
302     fprintf(stderr, "Missing required argument...");
303     fprintf(stderr, "%s", help);
304 }
305
306 initializeDCache(optionSetBits, optionAssociativity, optionBlockBits,
307                 optionVerboseMode);
308
309 FILE *fp = fopen(optionTraceFile, "r");
310 if (fp == NULL)
311 {
312     fprintf(stderr, "No such file or directory - %s", optionTraceFile);
313     return -1;
314 }
315
316 size_t max_length = 32;
317 char *line = (char *)malloc(max_length * sizeof(char));
318 while (getline(&line, &max_length, fp) != -1)
319 {
320     if (!isspace(line[0]))
321         continue;
322
323     char parsedCommand;
324     int parsedAddress;
325     int parsedBytes;
326     sscanf(line, " %c %x,%d", &parsedCommand, &parsedAddress,
327           &parsedBytes);
328     dCacheSimulate(parsedCommand, parsedAddress, parsedBytes);
329 }
330
331 fclose(fp);
332
333 printSummary(hits, misses, evictions);
334 return 0;
335 }

```

Listing 1: csim.c

The program was run against the reference simulator. The screenshot of the test benchmark run is presented on figure 1.

Points (s,E,b)	Your simulator			Reference simulator			
	Hits	Misses	Evicts	Hits	Misses	Evicts	
3 (1,1,1)	9	8	6	9	8	6	traces/yi2.trace
3 (4,2,4)	4	5	2	4	5	2	traces/yi.trace
3 (2,1,4)	2	3	1	2	3	1	traces/dave.trace
3 (2,1,3)	167	71	67	167	71	67	traces/trans.trace
3 (2,2,3)	201	37	29	201	37	29	traces/trans.trace
3 (2,4,3)	212	26	10	212	26	10	traces/trans.trace
3 (5,1,5)	231	7	0	231	7	0	traces/trans.trace
6 (5,1,5)	265189	21775	21743	265189	21775	21743	traces/long.trace
27							

TEST CSIM RESULTS=27

Figure 1: Benchmark results

## 2 Part B: Optimizing Matrix Transpose

The solution includes three different transpose functions for matrices with sizes:

- $32 \times 32$
- $64 \times 64$
- $32 \times 32$

The full source code of `trans.c` is show on listing 2.

```
1  #include <stdio.h>
2  #include "cachelab.h"
3
4  int is_transpose(int M, int N, int A[N][M], int B[M][N]);
5  void trans32(int M, int N, int A[N][M], int B[M][N]);
6  void trans64(int M, int N, int A[N][M], int B[M][N]);
7  void transX(int M, int N, int A[N][M], int B[M][N]);
8  char transpose_submit_desc[] = "Transpose submission";
9  void transpose_submit(int M, int N, int A[N][M], int B[M][N])
10 {
11     if (N == 32 && M == 32)
12     {
13         trans32(M, N, A, B);
14     }
15     else if (N == 64 && M == 64)
16     {
17         trans64(M, N, A, B);
18     }
19     else
20     {
21         transX(M, N, A, B);
22     }
23
24     return;
25 }
26
27 void transX(int M, int N, int A[N][M], int B[M][N])
28 {
29     for (int fourthI = 0; fourthI < N; fourthI += 8)
30     {
31         for (int fourthJ = 0; fourthJ < M; fourthJ += 8)
32         {
33             for (int sub = 0; sub < 8; sub += 4)
34             {
35                 int i = fourthI;
36                 int iEnd = i + 8 > N ? N : i + 8;
37                 for (; i < iEnd; i++)
38                 {
39                     int k = fourthJ + sub;
40                     int kEnd = k + 4 > M ? M : k + 4;
41                     for (; k < kEnd; k++)
42                     {
43                         B[k][i] = A[i][k];
44                     }
45                 }
46             }
47         }
48     }
49
50     return;
51 }
52
53 void trans64(int M, int N, int A[N][M], int B[M][N])
54 {
55     for (int fourthI = 0; fourthI < 64; fourthI += 8)
56     {
57         for (int fourthJ = 0; fourthJ < 64; fourthJ += 8)
58         {
59             for (int sub = 0; sub < 8; sub += 4)
60             {
61                 for (int i = 0; i < 8; i++)
62                 {
```

Listing 2 (Cont.): csim.c

```

63         int iEf = i + fourthI;
64         int a0 = A[iEf][fourthJ + sub];
65         int a1 = A[iEf][fourthJ + 1 + sub];
66         int a2 = A[iEf][fourthJ + 2 + sub];
67         int a3 = A[iEf][fourthJ + 3 + sub];
68         B[fourthJ + sub][iEf] = a0;
69         B[fourthJ + 1 + sub][iEf] = a1;
70         B[fourthJ + 2 + sub][iEf] = a2;
71         B[fourthJ + 3 + sub][iEf] = a3;
72     }
73 }
74 }
75 }
76
77     return;
78 }
79
80 void trans32(int M, int N, int A[N][M], int B[M][N])
81 {
82     for (int fourthI = 0; fourthI < 32; fourthI += 8)
83     {
84         for (int fourthJ = 0; fourthJ < 32; fourthJ += 8)
85         {
86             for (int i = 0; i < 8; i++)
87             {
88                 int iEf = i + fourthI;
89                 int a0 = A[iEf][fourthJ];
90                 int a1 = A[iEf][fourthJ + 1];
91                 int a2 = A[iEf][fourthJ + 2];
92                 int a3 = A[iEf][fourthJ + 3];
93                 int a4 = A[iEf][fourthJ + 4];
94                 int a5 = A[iEf][fourthJ + 5];
95                 int a6 = A[iEf][fourthJ + 6];
96                 int a7 = A[iEf][fourthJ + 7];
97                 B[fourthJ][iEf] = a0;
98                 B[fourthJ + 1][iEf] = a1;
99                 B[fourthJ + 2][iEf] = a2;
100                B[fourthJ + 3][iEf] = a3;
101                B[fourthJ + 4][iEf] = a4;
102                B[fourthJ + 5][iEf] = a5;
103                B[fourthJ + 6][iEf] = a6;
104                B[fourthJ + 7][iEf] = a7;
105            }
106        }
107    }
108
109    return;
110 }
111
112 /*
113  * registerFunctions - This function registers your transpose
114  * functions with the driver. At runtime, the driver will
115  * evaluate each of the registered functions and summarize their
116  * performance. This is a handy way to experiment with different
117  * transpose strategies.
118  */
119 void registerFunctions()
120 {
121     /* Register your solution function */
122     registerTransFunction(transpose_submit, transpose_submit_desc);
123 }
124
125 /*
126  * is_transpose - This helper function checks if B is the transpose of
127  * A. You can check the correctness of your transpose by calling
128  * it before returning from the transpose function.
129  */
130 int is_transpose(int M, int N, int A[N][M], int B[M][N])
131 {
132     int i, j;
133
134     for (i = 0; i < N; i++)
135     {

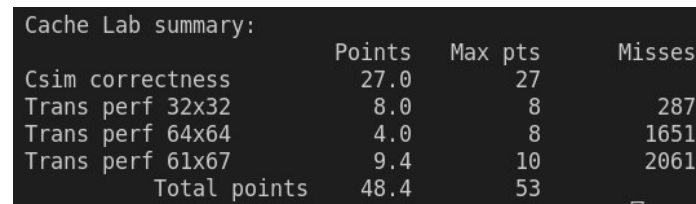
```

Listing 2 (Cont.): csim.c

```
136         for (j = 0; j < M; ++j)
137         {
138             if (A[i][j] != B[j][i])
139             {
140                 return 0;
141             }
142         }
143     }
144     return 1;
145 }
```

Listing 2: csim.c

The screenshot of the full lab benchmark run is presented on figure 2.



Cache Lab summary:			
	Points	Max pts	Misses
Csim correctness	27.0	27	
Trans perf 32x32	8.0	8	287
Trans perf 64x64	4.0	8	1651
Trans perf 61x67	9.4	10	2061
Total points	48.4	53	

Figure 2: Final lab results